

WHITEPAPER · 2026

Against AI Feudalism

How Ed4All Is Building the Missing Layer of the AI Economy

“Whoever owns the best-structured knowledge, not the largest model, will own the AI economy.”

Matthew D Murphy

INDEPENDENT RESEARCHER & CREATOR OF ED4ALL

Executive Summary

The numbers arrived before the story did. Enterprises have poured an estimated \$30–40 billion into generative AI, and — in preliminary findings — roughly 95% of organizations are getting zero measurable return (MIT NANDA, 2025). More than 80% of AI projects never reach meaningful production, about twice the failure rate of comparable work without an AI component, and the causes are overwhelmingly about data and organization, not model quality (RAND, 2024). I read those numbers not as a verdict on the technology but as an indictment of what sits beneath it: an infrastructure layer that was never built.

That missing layer is the argument of this paper. The constraint on AI value is not how large your model is; it is whether you can connect a model to what your institution already knows — in governed, portable, machine-readable form. Almost no one can. The knowledge exists, in syllabi and policy manuals and accreditation records, but it is trapped in formats no machine can use and owned, in practice, by no one. Institutions are data-rich and infrastructure-poor.

Into that vacuum flows the arrangement I call **AI feudalism**: you hand your documents to a platform, and from then on you rent access to a version of your own knowledge that you cannot export, inspect, or audit. No one ever decides to give it away. It leaves one routine upload at a time, until the asset you built is a line item on someone else's invoice.

Ed4All exists to stop that — to turn institutional documents into structured, validated, AI-ready knowledge the institution keeps, on open standards, with ownership and provenance intact at every step. The case is economic and political before it is technical.

Whoever owns the best-structured knowledge — not the largest model — will own the AI economy. For public institutions, whose knowledge was built with public money and ought to stay publicly governed, the question is not which vendor to pick. It is whether you will own the infrastructure of your own expertise

or rent it back forever. This paper makes that case with the evidence laid out plainly — every figure below links to its primary source — and argues that the choice should be made deliberately, and made now.

95%

of organizations get **zero measurable return** on generative-AI investment

MIT NANDA, 2025

80%+

of AI projects **fail to reach production** — about twice the non-AI IT rate

RAND, 2024

88%

organizational AI adoption — yet only **13% measure the return**

Stanford HAI · EDUCAUSE, 2026

The Problem No One Is Naming

Public discourse on artificial intelligence is organized around capability. Model releases are measured against one another; adoption is treated as a matter of when, not whether. The numbers support the optimism on that narrow axis: organizational AI adoption climbed from 78% in 2024 ([Stanford HAI, 2025](#)) to 88% a year later ([Stanford HAI, 2026](#)), and four in five university students now use generative AI ([Stanford HAI, 2026](#); [Digital Education Council, 2024](#)). On the question of whether the technology has arrived, the case is closed.

The question that has not been answered is whether it is working. The same period that produced near-universal adoption produced near-universal disappointment: in preliminary findings, 95% of organizations report zero measurable return on their generative-AI investment ([MIT NANDA, 2025](#)). More than 80% of AI projects never reach durable production — roughly twice the failure rate of IT projects without an AI component ([RAND, 2024](#)). The standard explanation — that the models are not yet good enough — does not survive the evidence. RAND's structured interviews with experienced practitioners locate the failure earlier in the pipeline: misunderstood problem definitions, inadequate data, and ungoverned content. Only one of its leading failure modes points primarily at the model itself ([RAND, 2024](#)).

The binding constraint is not model capability. It is **knowledge infrastructure** — the governed, machine-readable, reusable layer that must sit between what an institution already knows and the AI systems it wants to use. That layer is almost universally missing. Its absence is the problem no one is naming.

Data-rich, infrastructure-poor

Every institution this argument addresses is already sitting on the knowledge it needs. Universities hold course materials, learning outcomes, and accreditation documentation — mostly locked in PDFs and disconnected repositories. Government agencies hold policies, procedures, and compliance guidance scattered across shared drives and archived email. Employers hold operational know-how concentrated in the heads of their most experienced staff. In each case, the constraint is not the volume of content and not access to a model. It is the absence of a governed layer that makes institutional knowledge findable, citable, and reusable by a machine.

That diagnosis is not speculative. State technology leaders now rate data quality as a near-universal concern — 89% call it important — yet only 22% maintain a dedicated program to ensure it (NASCIO & EY, 2024). The appetite for AI-ready knowledge exists at the leadership level; the infrastructure to satisfy it does not. Generative AI has raised the stakes of this gap without closing it.

The hidden cost: institutional memory decay

The cost of missing knowledge infrastructure predates generative AI, and it is not small. The average interaction worker spent close to a fifth of the work-week simply searching for internal information even before the current era (McKinsey Global Institute, 2012); more than a quarter of a knowledge worker's time goes to locating what the organization already knows (McKinsey, 2012). When experienced staff leave, the tacit knowledge they carried goes with them — voluntary turnover alone costs U.S. employers roughly one trillion dollars annually (Gallup, 2019). New hires then spend months reconstructing institutional understanding that was never formally captured.

Generative AI does not retire this cost. It compounds it — introducing the additional risk of fluent, confident answers drawn from fragmented, outdated, or ungoverned material. There is a recurring pattern in how institutions manage knowledge under strain: elaborate systems can produce the appearance of understanding while the underlying comprehension quietly erodes. An AI system grounded in ungoverned content reproduces exactly that failure, and at machine speed. The problem is not the model's answer; it is that no one can trace the answer back to anything authoritative.

The real gap

The bottleneck is not what the models can do and not how much data an organization has. It is the absence of the governed, structured, machine-readable layer that transforms institutional documents into institutional knowledge — a layer that is findable, citable, provenance-tracked, and owned by the institution that created it.

Everything that follows turns on a single question: who builds that layer, and on whose terms.

What AI Feudalism Actually Looks Like

When an institution lacks knowledge infrastructure of its own, the path of least resistance is to hand its documents to an external platform and accept whatever arrangement comes with the service. The platform ingests the content, builds its own representations, and charges for access to those representations in perpetuity. The institution gets answers; the platform keeps the structure. What was an institutional asset — the accumulated know-how of departments, programs, and staff — silently migrates to a vendor's server, where it generates recurring revenue for the vendor rather than durable capital for the institution. This is the arrangement this paper calls **AI feudalism**: not because the technology is primitive, but because the property relation is. The people who work the land do not own it.

The structural literature on cloud computing has long documented how this outcome unfolds. Proprietary formats and interfaces are not incidental design choices; they are the mechanism by which lock-in is manufactured, and the costs become visible only when an organization tries to leave (Opara-Martins et al., 2016). Applied to institutional knowledge, that lock-in takes three distinct and compounding forms.

The first loss is portability. Knowledge uploaded to a vendor's platform is encoded in that vendor's proprietary representations — embedding formats, chunking schemas, retrieval indexes — that cannot be cleanly exported or transplanted. The institution can put its documents in; it cannot get its structured knowledge back out. Changing platforms means reconstructing the entire knowledge base from raw documents, an undertaking whose cost, once understood, tends to make incumbency permanent. Portability is not a convenience feature. It is the precondition for any meaningful vendor relationship.

The second loss is auditability. When retrieval is opaque — when there is no reliable mechanism to trace an AI-generated answer back to a specific, authoritative source — governance becomes impossible. For the domains where AI matters most in public institutions: compliance guidance, accreditation evidence, public-service delivery, legal and policy interpretation — an answer without a traceable source is a liability, not an asset. Retrieval opacity means an institution cannot answer the questions that regulators, auditors, and courts will eventually ask. It means accountability without understanding: the system can produce a confident output, but no one can say what it understood or what it got wrong.

The third loss is ownership in any meaningful sense. Responsible vendors have moved quickly to disclaim training on enterprise content. Microsoft states that data accessed through its enterprise graph "aren't used to train foundation models" (Microsoft, 2025); OpenAI says that "by default, we do not train on any inputs or outputs ... for business users" (OpenAI, 2025); GitHub similarly exempts enterprise and organization data from training use (GitHub, 2026). These commitments matter — but they are defaults set unilaterally by the vendor, varying by contract tier and subject to revision at the vendor's discretion. They are not guarantees anchored in the institution's own control of its data. An institution whose knowledge governance depends on the current language of a counterparty's policy page does not own its knowledge; it is merely borrowing it back under favorable terms.

The deeper problem is one of institutional recognition. There is a pattern, familiar across sectors, by which capable organizations quietly cease to acknowledge value they worked to build — where hard-earned knowledge, rigorously documented expertise, and validated practice gradually lapse into dependency on external systems, not through any single deliberate decision but through a sequence of individually reasonable steps, each of which leaves every formal commitment nominally intact. AI feudalism is the same pattern indexed to the current moment: an institution's knowledge is absorbed into a platform, abstracted into representations the institution does not hold, and rented back at a subscription rate that rises as the alternative recedes. No one decided to transfer the asset. It transferred through inaction.

Avoiding that outcome is not primarily a matter of vendor selection or contract negotiation, though both matter. It is a matter of building the infrastructure layer that makes ownership real before uploading the first document — the layer that preserves portability, enforces auditability, and grounds every AI interaction in knowledge the institution actually holds and can inspect. That is the structural argument for what follows.

Ed4All: Knowledge Infrastructure, Not a Chatbot

Ed4All is not a model and not a chatbot. It is a five-layer pipeline that converts institutional documents into structured, governed, AI-ready knowledge repositories the institution owns outright. The right frame for understanding each layer is not what it does mechanically but what institutional value it preserves or creates — and why the absence of that layer is a structural liability rather than a mere inconvenience. Every layer is built on open, inspectable standards, so that no single vendor sits as an intermediary between the institution and its own knowledge.

Layer 1: Accessibility remediation — the right to use without surrendering anything. The pipeline begins by converting legacy documents into content that conforms to WCAG 2.1 Level AA and Section 508. The economic logic here is compact: accessibility and machine-readability are not separate projects but two out-

puts of a single act of structuring. A document that cannot be parsed by a screen reader also cannot be parsed reliably by a language model. Making knowledge usable for a person with a disability and making it parseable by an AI system require the same transformation — semantic structure, well-formed headings, descriptive metadata — applied once. There is a broader principle at stake that extends beyond efficiency: access to institutional knowledge should not require a person — or an institution — to disclose or surrender something about themselves as the price of admission. Remediation at this layer is not a compliance burden; it is the act of asserting that the knowledge belongs to those who created it and to those it was created to serve. The approaching Title II deadlines — April 2027 for larger entities, 2028 for smaller ones — mean this transformation is already required by law (DOJ, 2024; DOJ, 2026). Building AI-readiness on top of a mandate already in motion is not a choice; it is an available efficiency that most institutions are currently leaving on the table.

Layer 2: Semantic chunking with provenance — from pages to citable knowledge units. Once documents are structurally sound, content is segmented by meaning rather than by arbitrary length, and each segment is bound to a precise provenance record: the source document, the section, the version, and the date of record. This is the transformation that turns a PDF library into a set of retrievable, citable knowledge units — the difference between "the answer is somewhere in these files" and "the answer is this passage, in this document, at this version, and here is the chain of custody." Provenance is not metadata decoration; it is the precondition for auditability. A knowledge unit that cannot be traced to its source cannot be governed, and ungoverned units are how plausible-sounding wrong answers propagate at scale. Every subsequent layer depends on this one.

Layer 3: Knowledge graph on open standards — making implicit relationships explicit and portable. Concepts, outcomes, prerequisites, and the relationships among them are represented as a knowledge graph using open semantic-web standards: JSON-LD, RDF, SHACL, and OWL (W3C), with course-packaging and competency structures expressed in open interchange formats (1EdTech). The institutional value this layer creates is substantial. Relationships that previously lived in an expert's head — this outcome is a prerequisite for that one; this policy clause qualifies that procedure — become explicit, queryable, and reusable. Curriculum alignment, gap analysis, and compliance mapping that are currently done by hand become computable. And because the representation is open and standardized, any compliant tool, model, or system can operate over the

same knowledge base. The institution is not locked into any vendor's interpretation of what its knowledge means; it retains the full ability to migrate, audit, extend, or share the graph on its own terms.

Layer 4: Validation and governance — the difference between accountability and understanding. Schema constraints enforce structural integrity across the repository, and behavioral evaluation probes test the system's responses against the underlying content — including, critically, the discipline of abstaining when the repository does not support a given answer. This layer is what converts a knowledge repository from a latent liability into an auditable asset. It is also where the deepest question in this paper becomes operational. A recurring theme in this author's work is the distance between accountability and genuine understanding: the capacity to build elaborate machinery that produces the appearance of rigor while the underlying comprehension is absent. An AI system grounded in ungoverned content is the same failure replicated at machine speed — it has not understood anything; it has learned to sound as though it has. Governed knowledge infrastructure is designed to refuse exactly that. The NIST AI Risk Management Framework's core functions — GOVERN, MAP, MEASURE, MANAGE (NIST, 2023) and its generative-AI profile (NIST, 2024) — describe precisely the lifecycle a repository built at this layer is designed to support. Governance is not a feature to be added later; it is a structural property of the repository itself or it is nothing at all.

Layer 5: Grounded retrieval and local model adaptation — local capability without structural dependency. The governed, provenance-bound, validated graph supports hybrid retrieval indexes for grounded question-answering: answers are drawn from owned, cited content rather than from a model's parametric memory. Where an institution wants deeper adaptation, small models can be fine-tuned on the institution's own validated knowledge through adapter training — without transmitting that knowledge to an outside platform as the price of the improvement. The result is local AI capability grounded in governed, owned content. The institution gets the benefits of model-level adaptation without accepting the lock-in that typically accompanies it.

The pipeline, read as institutional value — each layer preserves something the institution owns.

- 1 **Accessibility remediation.** Legacy documents become standards-conformant, machine-parseable content (WCAG 2.1 AA, Section 508) — human-usable and machine-readable in one act.
- 2 **Semantic chunking with provenance.** Content is segmented by meaning and bound to its source location — pages become citable knowledge units.
- 3 **Knowledge graph on open standards.** Concepts, outcomes, and relationships represented in JSON-LD / RDF / SHACL / OWL — portable and inspectable.
- 4 **Validation & governance.** Schema constraints and evaluation probes enforce integrity, including abstaining when an answer is unsupported.
- 5 **Grounded retrieval & local adaptation.** Hybrid retrieval and optional small-model adaptation deliver owned, governed capability — no structural vendor dependency.

A worked example — what a single ingestion produces.

INPUT A community college exports **200 course syllabi** and their supporting documents into the pipeline.



- ▶ **Accessible HTML** for every document (WCAG 2.1 AA / Section 508) — usable by a student with a disability and parseable by a machine in the same pass.
- ▶ **A JSON-LD / RDF knowledge graph** linking courses, learning outcomes, prerequisites, and competencies.
- ▶ **A SHACL validation report** flagging missing outcomes, broken prerequisite chains, and unmapped competencies before they reach a catalog.
- ▶ **Provenance-bound chunks**, each traceable to its source document, section, and version.
- ▶ **A local retrieval tutor** answering students and advisors from the college's own validated content, with citations — and abstaining when the content does not support an answer.

One ingestion. The college owns every artifact, in open formats, and reuses them across advising, accreditation, articulation, and program review — without redoing the work.

Open standards as an economic choice

The consistent use of open formats throughout this pipeline — HTML, JSON-LD, RDF, SHACL, OWL, open course-packaging and competency-exchange standards (W3C; 1EdTech), and accessibility standards (WCAG 2.1 Level AA; Section 508) — is not a technical preference. It is an economic and governance decision with compounding consequences.

Proprietary formats — whether in vector databases, embedding stores, or knowledge-graph representations — accomplish something specific: they bind the institution's knowledge to the vendor's system. The institution can put its documents in; it generally cannot get its structured knowledge back out in a form another system can use. The structural literature on cloud computing documents exactly this mechanism: proprietary interfaces are the primary vehicle through which lock-in is established, and the cost is typically invisible until the institution attempts to leave (Opara-Martins et al., 2016). Applied to knowledge infrastructure, the cost is compounded: the institution loses not merely a software subscription but the accumulated structure — the graph, the provenance chains, the validated relationships — that represents years of organizational effort.

Open standards dissolve that bind. When a knowledge repository is expressed in JSON-LD, RDF, and SHACL, it is readable by any W3C-compliant tool. When competency frameworks are expressed in CASE and course packages in Common Cartridge, they are portable across any 1EdTech-compliant system (1EdTech). The institution's choice of downstream model, retrieval system, or platform vendor becomes a procurement decision rather than an existential dependency. Equally important, the knowledge can be inspected: a SHACL-validated RDF graph can be queried, audited, and extended by anyone with standard tooling, with no proprietary interpreter required.

The format question is where ownership is either retained or quietly signed away — often without any single decision that feels momentous. Choosing open standards at every layer is the architectural commitment that makes all other ownership claims credible.

Three ways to connect an institution to AI — only one preserves ownership.

PROPERTY	VECTOR DATABASE	KNOWLEDGE GRAPH	GOVERNED REPOSITORY
Portability of structured knowledge	✗	<i>partial</i>	✓
Provenance to source & version	✗	<i>partial</i>	✓
Auditable, traceable answers	✗	<i>partial</i>	✓
Open, inspectable formats	✗	✓	✓
Institution retains ownership	✗	<i>partial</i>	✓

Three Sectors Where This Changes the Math

The promise of knowledge infrastructure is only worth making if it cashes out in specific, observable terms. In the three sectors that follow, the problem is the same structural one — institutional knowledge exists but cannot be used — but the economic injury is distinct in each case, and so is the payoff when the structure changes.

Higher education: curriculum as reusable capital

Every course revision cycle, the same intellectual work is done again from scratch: outcomes drafted, content sequenced, assessment criteria defined. The results live in syllabi and learning management systems as human-readable prose, and when a course is archived or an instructor moves on, the work is largely gone. Alignment — demonstrating that outcomes match assessments, that prerequisites are coherent, that accreditation standards are met — is done by hand, repeatedly, across programs that share substantial underlying content.

Expressing curriculum as a graph of outcomes, concepts, prerequisites, and assessments converts that perishable effort into durable infrastructure. A single structured representation of a course can simultaneously support AI-assisted tu-

toring, competency mapping, gap analysis, adaptive sequencing, and accreditation documentation — without redundant rebuilds. The initial structuring work does not disappear at the end of the semester; it compounds.

The urgency is real from both directions. Students have already made the transition: 86% incorporate AI into their studies, and more than half do so weekly ([Digital Education Council, 2024](#)). Institutions have not kept pace. Only 13% are even measuring the return on their AI investments ([EDUCAUSE, 2026](#)), which means the dominant response to widespread student adoption has been neither resistance nor intentional integration but ambiguity. Meanwhile, instructors who use AI daily report a 36% workload decrease ([Tyton Partners & D2L, 2025](#)) — a signal that the value is there when the tools are grounded in reliable content.

Structured curriculum is how that reliability is built. The GI Bill generation understood that public investment in education was not recurring cost but compounding capital — it returned to the economy in workforce capability, tax revenue, and institutional knowledge for decades. The same logic applies here: treating curriculum as durable, reusable infrastructure is not a technical choice but a decision about what kind of investment education is.

Government: institutional memory that outlasts its holders

Government agencies run on expertise. Policy decisions, procurement judgments, compliance determinations, and service delivery workflows are inseparable from the knowledge held by the staff who have spent years building it. That knowledge is rarely written down in any usable form. When experienced personnel retire or move on, the continuity of service degrades — not dramatically, but cumulatively, in ways that compound over time and surface only when something goes wrong.

The cost of this is not hypothetical. When policy manuals, workflow documentation, legal guidance, and training materials are converted into a validated, queryable knowledge repository, a new employee can ask a specific procedural question and receive an answer traceable to its authoritative source document, at its current version, with a clear audit trail. The answer is not produced by a model reasoning from general knowledge; it is retrieved from the agency's own governed content, with full provenance. Compliance evidence is generated as a natural byproduct of normal operations rather than assembled under deadline.

The diagnosis is one that state technology leaders already share clearly: 89% consider data quality important, yet only 22% maintain a dedicated program to ensure it (NASCIO & EY, 2024). The gap is not skepticism about the value of structured data — it is the absence of a practical path from the current state to governed infrastructure. Generative AI is rising on the state technology agenda (NASCIO, 2024), and state governments are actively working through what it means for IT workforces and service delivery (NASCIO & McKinsey, 2024). But enthusiasm for AI tools does not close the data-quality gap; it widens the cost of leaving it unaddressed. Deploying AI over ungoverned content means faster answers with less traceability — a combination that is worse, not better, for regulated public services.

Governed knowledge infrastructure does not make AI safe at the level of policy — that is a governance and accountability question beyond any technical system. But it does make AI auditable at the level of content: every answer can be traced to the document that authorized it, and every document can be validated, versioned, and owned by the institution that created it.

Workforce development: closing the gap between credentials and competencies

Career and technical education programs are designed around labor market outcomes, but the connection between curriculum and employer demand is usually indirect. Programs are built from state standards and faculty expertise, then adjusted based on employer feedback that arrives slowly, informally, and unevenly. By the time a misalignment between what a program teaches and what regional employers require becomes visible enough to act on, students have already completed credentials that do not fully match the market.

The scale of the problem is documented. Nationally, 28% of middle-skills credentials have no direct occupational match; in roughly half of local labor markets, at least half of all middle-skills credentials would need to be granted in different fields to align with projected demand (Georgetown CEW, 2024). This is not a marginal inefficiency. It is a structural failure of signal transmission between labor markets and the programs that train for them — and it is happening against a backdrop of accelerating skill change, as the required content of jobs shifts faster than curriculum cycles typically allow (Lightcast, 2024).

When curricula are expressed as machine-readable competency graphs, the alignment problem becomes structural rather than anecdotal. Regional colleges, workforce boards, and employers can map their competency nodes to the same vocabulary, making gaps explicitly visible — not as impressions gathered in advisory meetings, but as measurable distances between what a program certifies and what employers specify. Misalignments that currently require a cycle of feedback, revision, and reapproval to fix can instead be identified and addressed at the granular level of specific competency nodes.

This is also a direct fulfillment of federal statutory intent. The Carl D. Perkins Career and Technical Education Act, funded at \$1.44 billion for the FY 2024 Basic State Grant, calls for a "nonduplicative sequence of academic and technical content" leading to recognized credentials ([ACTE, 2025](#)). That language is aspirational in the absence of infrastructure that makes sequencing legible and verifiable. Structured, machine-readable curriculum is how a statutory aspiration becomes an operational standard — turning Perkins' coherence requirement from a compliance claim into a measurable, improvable property of the program itself.

What Governed Knowledge Infrastructure Requires

Building a knowledge repository is not primarily a technical undertaking; it is a governance undertaking. Three properties are non-negotiable, and they must be engineered in from the start — they cannot be patched on after deployment.

Provenance. Every knowledge unit must carry a complete chain of custody: source document, section identifier, version, and ingestion timestamp. This is not a record-keeping preference; it is the condition under which an answer can be traced to an authorizing document. In regulated contexts — compliance guidance, credentialing, public services — an answer that cannot be anchored to its source is not an asset. It is a liability with a latency period.

Validation. Structure and content must be held to enforceable constraints. Schema-level rules (expressed in open standards such as SHACL) catch structural malformation before it reaches retrieval. Behavioral probes — queries whose cor-

rect answer is known — test whether the system draws the right boundaries: does it answer accurately when the content supports an answer, and does it abstain when it does not? A system that cannot abstain has not learned to reason from evidence; it has learned to fill silence with plausible text. There is a meaningful difference between accountability — the ability to produce an output — and understanding — the capacity to show why that output is warranted. Ungoverned AI conflates the two; governed infrastructure is designed to maintain the distinction.

Auditability. Every change to the repository — content added, revised, flagged, or retired — must be logged. For institutions operating in accountability environments, this is not optional: regulators, accreditors, and courts do not accept "the model said so" as a justification. Auditability is what converts an AI system from a black box that produces conclusions into an instrument that produces auditable decisions.

These three properties map directly onto the vocabulary that federal guidance has adopted for trustworthy AI. The NIST AI Risk Management Framework organizes AI governance into four interacting functions: GOVERN establishes organizational policies, roles, and accountability structures; MAP identifies the context and risks associated with an AI deployment; MEASURE quantifies those risks with appropriate metrics; and MANAGE applies that measurement to ongoing operations, remediation, and monitoring (NIST, 2023). The companion Generative AI Profile extends this framework to the specific risk surface of large language models — including content provenance, confabulation, and data quality failures (NIST, 2024). A governed knowledge repository is the artifact that makes GOVERN, MAP, MEASURE, and MANAGE operational rather than aspirational: without provenance-linked, validated, auditable content, the framework has nothing to act on.

Why ingestion is not infrastructure

The most common error in AI deployment is treating document ingestion as knowledge infrastructure. It is not. Uploading documents to a vector database produces a faster retrieval surface over ungoverned content. That is a different thing — and, at scale, a more dangerous thing — than a governed knowledge repository.

The failure modes that turn AI from productivity multiplier to institutional liability are not exotic. They are predictable from first principles, and they are the direct consequence of confusing ingestion with infrastructure:

Provenance collapse. When retrieved content is not linked to its source, confident-sounding answers cannot be traced. An output that cannot be verified cannot be acted on responsibly. The system has manufactured the appearance of knowledge without any mechanism for checking it.

Version drift. Policies change. Regulations are revised. Procedures are updated. An ingestion-only pipeline that re-indexes periodically without version control will silently serve outdated guidance alongside current guidance, with no signal to distinguish them. Institutions that rely on such systems for compliance-sensitive decisions are not managing risk; they are deferring its discovery.

Bias amplification at machine scale. Source documents encode the assumptions, omissions, and language of the moment they were written. A retrieval system without bias-detection and flagging capabilities reproduces those assumptions at the speed of query. What was a slow-moving institutional bias becomes a fast-moving operational one.

Confabulation with institutional authority. Generative AI systems produce fluent text even when the underlying retrieval is thin or misdirected. When that text is presented as the output of an institutional knowledge system, it carries implied authority. The NIST Generative AI Profile specifically identifies confabulation as a primary risk category requiring active MEASURE and MANAGE controls (NIST, 2024). An ingestion-only pipeline has no mechanism to apply them.

Vendor lock-in of the knowledge layer. When retrieval depends on proprietary embedding formats or platform-specific index structures, the institution cannot export, inspect, or migrate its structured knowledge. The content is in; the knowledge cannot get out.

Access without accountability. If the repository does not enforce role-based access and log retrieval events, the institution cannot demonstrate that sensitive content was accessed appropriately — a requirement in virtually every regulated sector.

Each failure mode is an argument for governed infrastructure, not an argument against AI. The institutions that will extract durable value from these systems are the ones that treat governance as a prerequisite rather than a retrofit — building the provenance, validation, and auditability layers before deployment, not in response to an incident after it.

The Policy Window

The strongest near-term forcing function for public AI infrastructure is not a new appropriation or a voluntary commitment. It is a compliance deadline that is already on the calendar.

The Department of Justice's Title II final rule requires state and local government web content to meet WCAG 2.1 Level AA (DOJ, 2024). Compliance dates — April 26, 2027 for entities serving populations of 50,000 or more, and 2028 for smaller jurisdictions (DOJ, 2026) — are close enough that procurement cycles for meeting them are opening now. What the mandate forces, technically, is the conversion of legacy institutional content into structured, semantically coherent, machine-readable form. That transformation is indistinguishable from the first step in building AI-ready knowledge infrastructure. An agency or institution that meets Title II well — not merely as a check-the-box pass but as a genuine restructuring of its content — will simultaneously have produced a governed, accessible knowledge base that any compliant AI system can operate over. The compliance deadline is, in effect, a knowledge-infrastructure deadline. Policy should treat it as one.

That framing is reinforced by the current federal procurement posture. The operative AI-acquisition memorandum directs agencies to scope intellectual-property rights "to avoid vendor lock-in" and to preserve the components needed to operate and monitor a procured system within the acquiring agency's own control (OMB, 2025) — language that strengthens the anti-lock-in posture of its predecessor (OMB, 2024). Federal data-rights provisions give agencies the contractual tools to negotiate ownership and interoperability terms for any knowledge generated or structured under a government contract (FAR Part 27). The legal and

regulatory infrastructure for an ownership-first procurement posture already exists. What is missing is consistent application of it to AI-knowledge contracts specifically. Three actionable directions follow.

Procurement language that preserves public ownership. Every public AI contract involving institutional content should include three non-negotiable terms: first, an exportability requirement obligating the vendor to return all ingested and structured content in open, portable formats that the institution controls without ongoing vendor access; second, an auditability requirement mandating that AI outputs be traceable to source documents with version and provenance metadata retained by the institution, not solely by the platform; and third, a use restriction prohibiting the vendor from using publicly funded institutional knowledge for model training or product improvement without explicit, revocable, written consent. These terms do not conflict with the current executive guidance on AI (EO 14179, 2025); they operationalize its principle that government should be AI-capable without becoming AI-dependent. Section 508 of the Rehabilitation Act already imposes analogous structural accountability for accessibility in federal IT (Section 508); the same logic applies to knowledge ownership.

Open-standards mandates. Public procurement of AI-knowledge systems should require that all structured outputs — semantic representations, competency maps, curriculum graphs, policy repositories — be expressed in open, inspectable, portable formats: HTML for accessible content; JSON-LD, RDF, and SHACL for semantic structure (W3C); open course-packaging and competency-exchange standards (1EdTech) for education-sector deployments. Requiring open formats is not a technical preference; it is the structural guarantee that the institution retains the ability to migrate, audit, extend, or share its knowledge base without vendor permission. Proprietary embedding formats and closed graph schemas do the opposite: they encode the institution's knowledge in a representation only the vendor can read, making every future procurement decision a hostage negotiation. WCAG 2.1 Level AA and Section 508 compliance should be treated as pre-conditions for AI-readiness procurement eligibility, not as separate line items — reinforcing the Title II alignment described above and ensuring that accessibility and machine-usability are achieved in one pass rather than two.

Regional knowledge-infrastructure investment. Community colleges, regional universities, and workforce boards are the natural institutional anchors for locally governed knowledge infrastructure — they hold the curriculum, competency

frameworks, and employer relationships that make regional labor markets legible. Several federal vehicles are already aligned to this function. The Carl D. Perkins Act — the primary federal vehicle for career-and-technical education, funded at \$1.44 billion for the FY 2024 Basic State Grant — explicitly calls for a "nonduplicative sequence of academic and technical content" leading to recognized credentials (ACTE, 2025); structured, machine-readable curriculum expressed in open standards is precisely how that statutory language becomes operational rather than aspirational. WIOA's nondiscrimination framework governs workforce system access and equity (29 CFR Part 38), and its logic extends naturally to the accessibility of AI-mediated career services. The National AI Research Resource (NSF, 2025) represents an emerging national shared-infrastructure model whose architecture could inform analogous regional efforts. Treating governed knowledge infrastructure as a regional competitiveness metric — alongside broadband access, cloud capacity, and technical workforce supply — would direct public investment toward durable local capacity rather than recurring vendor subscriptions. The thesis that knowledge, once publicly funded, should remain publicly governed is not a political assertion but an economic one: publicly funded knowledge that migrates into proprietary representations is a transfer of public capital to private balance sheets, at scale, with no appropriations vote.

None of the three directions requires new statute to begin. The procurement tools, the data-rights framework, the accessibility mandates, and the workforce-investment vehicles are already in place. What is required is the deliberate, coordinated use of them — treating knowledge ownership as a procurement objective rather than an afterthought, and the approaching Title II deadlines as the clearest near-term opportunity to do both at once.

The Stake

The advantage in AI-enabled markets will not go to whoever rents the largest model. It will go to whoever owns the best-structured knowledge the model stands on. Build that infrastructure and you keep the long-term value of your own expertise. Fail to, and you will buy it back, year after year, at a price someone else sets.

The stakes are sharpest right now. Institutions are operating under real strain — eroded public trust, deep polarization, contested authority — and in that climate, control of institutional knowledge is not a background condition. It is a live variable. The default path is not neutral: it quietly concentrates the value of public and institutional knowledge in a handful of platform owners, exactly when the question of who speaks with institutional authority is already in dispute.

The future of AI adoption depends less on who owns the largest model and more on who owns the best-structured knowledge.

Ed4All is the infrastructure that keeps that ownership local. The alternative is AI feudalism — not a dramatic rupture, but a quiet succession of unremarkable decisions that hand an institution's hard-won knowledge to a subscription it can neither export nor audit. The choice between those futures is not a technical one. It is economic, institutional, and — for publicly funded knowledge above all — political. It is also still ours to make.

Ed4All is open-source. The pipeline, schemas, validation profiles, and tooling are available at github.com/mdmurphy822/Ed4All.

References

Every source is publicly available at the link provided, and each cited figure was checked against that primary source — follow any link to confirm a claim directly. Links resolve to the canonical publisher or government document, not a local copy. Sources noted as "preliminary" (notably the 2025 MIT NANDA findings) are early, fast-moving results and should be read as such; figures available only in paywalled reports were omitted by design.

1EdTech. Common Cartridge; CASE (Competencies and Academic Standards Exchange). <https://www.1edtech.org/standards/cc>

- ACTE. (2025). Introduction to the Carl D. Perkins CTE Act. <https://www.acteonline.org/wp-content/uploads/2025/03/ACTE-Perkins101-Basics-Feb2025.pdf>
- Digital Education Council. (2024). Global AI Student Survey 2024. <https://www.digitaleducationcouncil.com/dec-insights/how-students-use-ai-the-evolving-relationship-between-ai-and-higher-education>
- DOJ. (2024). ADA Title II Final Rule, 89 Fed. Reg. 31320. <https://www.govinfo.gov/content/pkg/FR-2024-04-24/pdf/2024-07758.pdf>
- DOJ. (2026). Interim Final Rule: Extension of ADA Title II Web Accessibility Compliance Dates. <https://www.federalregister.gov/documents/2026/04/20/2026-07663/extension-of-compliance-dates-for-nondiscrimination-on-the-basis-of-disability-accessibility-of-web>
- EDUCAUSE. (2026). The Impact of AI on Work in Higher Education. <https://www.educause.edu/research/2026/the-impact-of-ai-on-work-in-higher-education>
- Executive Order 14179. (2025). Removing Barriers to American Leadership in Artificial Intelligence. <https://www.govinfo.gov/content/pkg/FR-2025-01-31/pdf/2025-02172.pdf>
- FAR Part 27. Patents, Data, and Copyrights. <https://www.acquisition.gov/far/part-27>
- Gallup. (2019). This Fixable Problem Costs U.S. Businesses \$1 Trillion. <https://www.gallup.com/workplace/247391/fixable-problem-costs-businesses-trillion.aspx>
- Georgetown CEW. (2024). The Great Misalignment. <https://cew.georgetown.edu/cew-reports/greatmisalignment/>
- GitHub. (2026). Updates to Our Privacy Statement and Terms of Service: How We Use Your Data. <https://github.blog/changelog/2026-03-25-updates-to-our-privacy-statement-and-terms-of-service-how-we-use-your-data/>
- Lightcast. (2024). The Speed of Skill Change (US edition). https://4906807.fs1.hubspotusercontent-na1.net/hubfs/4906807/Speed%20of%20Skill%20Change_US_Lightcast-final.pdf
- McKinsey. (2012). Rethinking Knowledge Work: A Strategic Approach. <https://www.mckinsey.com/capabilities/people-and-organizational-performance/our-insights/rethinking-knowledge-work-a-strategic-approach>
- McKinsey Global Institute. (2012). The Social Economy: Unlocking Value and Productivity Through Social Technologies. <https://www.mckinsey.com/industries/technology-media-and-telecommunications/our-insights/the-social-economy>

- Microsoft. (2025). Enterprise Data Protection in Microsoft 365 Copilot. <https://learn.microsoft.com/en-us/microsoft-365/copilot/enterprise-data-protection>
- MIT NANDA. (2025). The GenAI Divide: State of AI in Business 2025 (v0.1 preliminary). https://mlq.ai/media/quarterly_decks/v0.1_State_of_AI_in_Business_2025_Report.pdf
- NASCIO. (2024). 2024 State CIO Survey: Building Blocks of the Next Generation CIO. https://www.nascio.org/wp-content/uploads/2024/09/NASCIO_2024-State-CIO-Survey_a11y-1.pdf
- NASCIO & EY. (2024). Is Your State Data Quality Ready for GenAI? https://www.nascio.org/wp-content/uploads/2024/09/NASCIO_EY_Data-Quality_2024.pdf
- NASCIO & McKinsey. (2024). Generative AI and Its Impact on State Government IT Workforces. https://www.nascio.org/wp-content/uploads/2024/04/NASCIO_McKinsey_GenAI-and-Impact-on-State-IT-Workforces_2024.pdf
- NIST. (2023). AI Risk Management Framework 1.0 (NIST AI 100-1). <https://doi.org/10.6028/NIST.AI.100-1>
- NIST. (2024). AI RMF: Generative AI Profile (NIST AI 600-1). <https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.600-1.pdf>
- NSF. (2025). NSF Announces Funding to Establish National AI Research Resource (NAIRR) Operations Center. <https://www.nsf.gov/news/nsf-announces-funding-establish-national-ai-research>
- OMB. (2024). M-24-18: Advancing the Responsible Acquisition of AI in Government (rescinded predecessor). <https://www.whitehouse.gov/wp-content/uploads/2024/10/M-24-18-AI-Acquisition-Memorandum.pdf>
- OMB. (2025). M-25-22: Driving Efficient Acquisition of Artificial Intelligence in Government. <https://www.whitehouse.gov/wp-content/uploads/2025/02/M-25-22-Driving-Efficient-Acquisition-of-Artificial-Intelligence-in-Government.pdf>
- OpenAI. (2025). How Your Data Is Used to Improve Model Performance. <https://openai.com/policies/how-your-data-is-used-to-improve-model-performance/>
- Opara-Martins, J., Sahandi, R., & Tian, F. (2016). Critical Analysis of Vendor Lock-in and Its Impact on Cloud Computing Migration: A Business Perspective. *Journal of Cloud Computing*. <https://link.springer.com/content/pdf/10.1186/s13677-016-0054-z.pdf>
- RAND. (2024). The Root Causes of Failure for Artificial Intelligence Projects and How They Can Succeed. https://www.rand.org/content/dam/rand/pubs/research_reports/RRA2600/RRA2680-1/RAND_RRA2680-1.pdf

- Section 508 of the Rehabilitation Act (29 U.S.C. 794d). <https://www.section508.gov/manage/laws-and-policies/section-508-law/>
- Stanford HAI. (2025). 2025 AI Index Report. https://hai.stanford.edu/assets/files/hai_ai_index_report_2025.pdf
- Stanford HAI. (2026). 2026 AI Index Report. https://hai.stanford.edu/assets/files/ai_index_report_2026.pdf
- Tyton Partners & D2L. (2025). Time for Class 2025. https://www.d2l.com/newsroom/tyton_partners_report_examines_ai_in_higher_education/
- W3C. JSON-LD 1.1, RDF, SHACL, SKOS (Recommendations). <https://www.w3.org/TR/json-ld11/>
- WIOA Nondiscrimination, 29 CFR Part 38. <https://www.ecfr.gov/current/title-29/subtitle-A/part-38>